

Introduction to Hadoop and HDFS

Lecture 2

October 19th, 2017

Jae W. Lee (jaewlee@snu.ac.kr)

Computer Science and Engineering

Seoul National University

Slide credits: Rafael Coss (Hortonworks); Owen O'Malley (Hortonworks);
DarrelAucoin (Intro to Hadoop); Tom White (Hadoop, The Definitive Guide)

Outline

- **What is Apache Hadoop?**
- Key Components of Hadoop Stack
- Hadoop Distributed File System (HDFS)
- HDFS File System Shell Commands
- Apache Spark on HDFS

What is Apache Hadoop?

■ Solution for Big Data

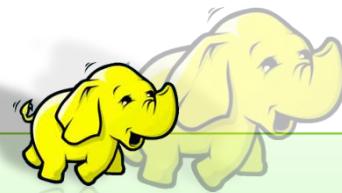
- Designed for volume, velocity, variety & complexity of data

■ Data Platform Deployed on Commodity Hardware that

- Stores petabytes of data reliably
- Runs huge distributed applications
- Enables a rational economics model

■ Set of Open Source Projects

- Apache Software Foundation
- Loosely coupled, ship early/often

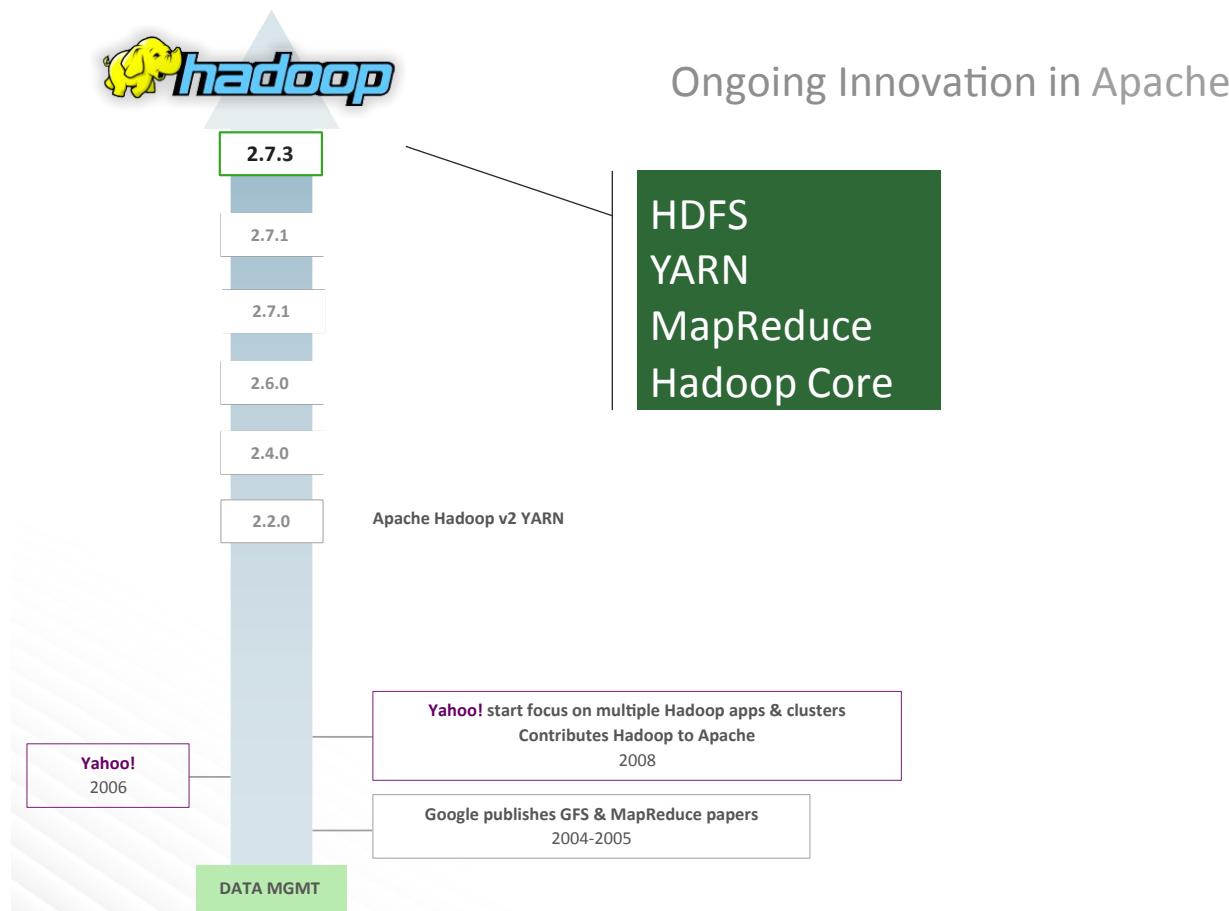


One of the best examples of open source driving innovation and creating a market

Source: Hortonworks

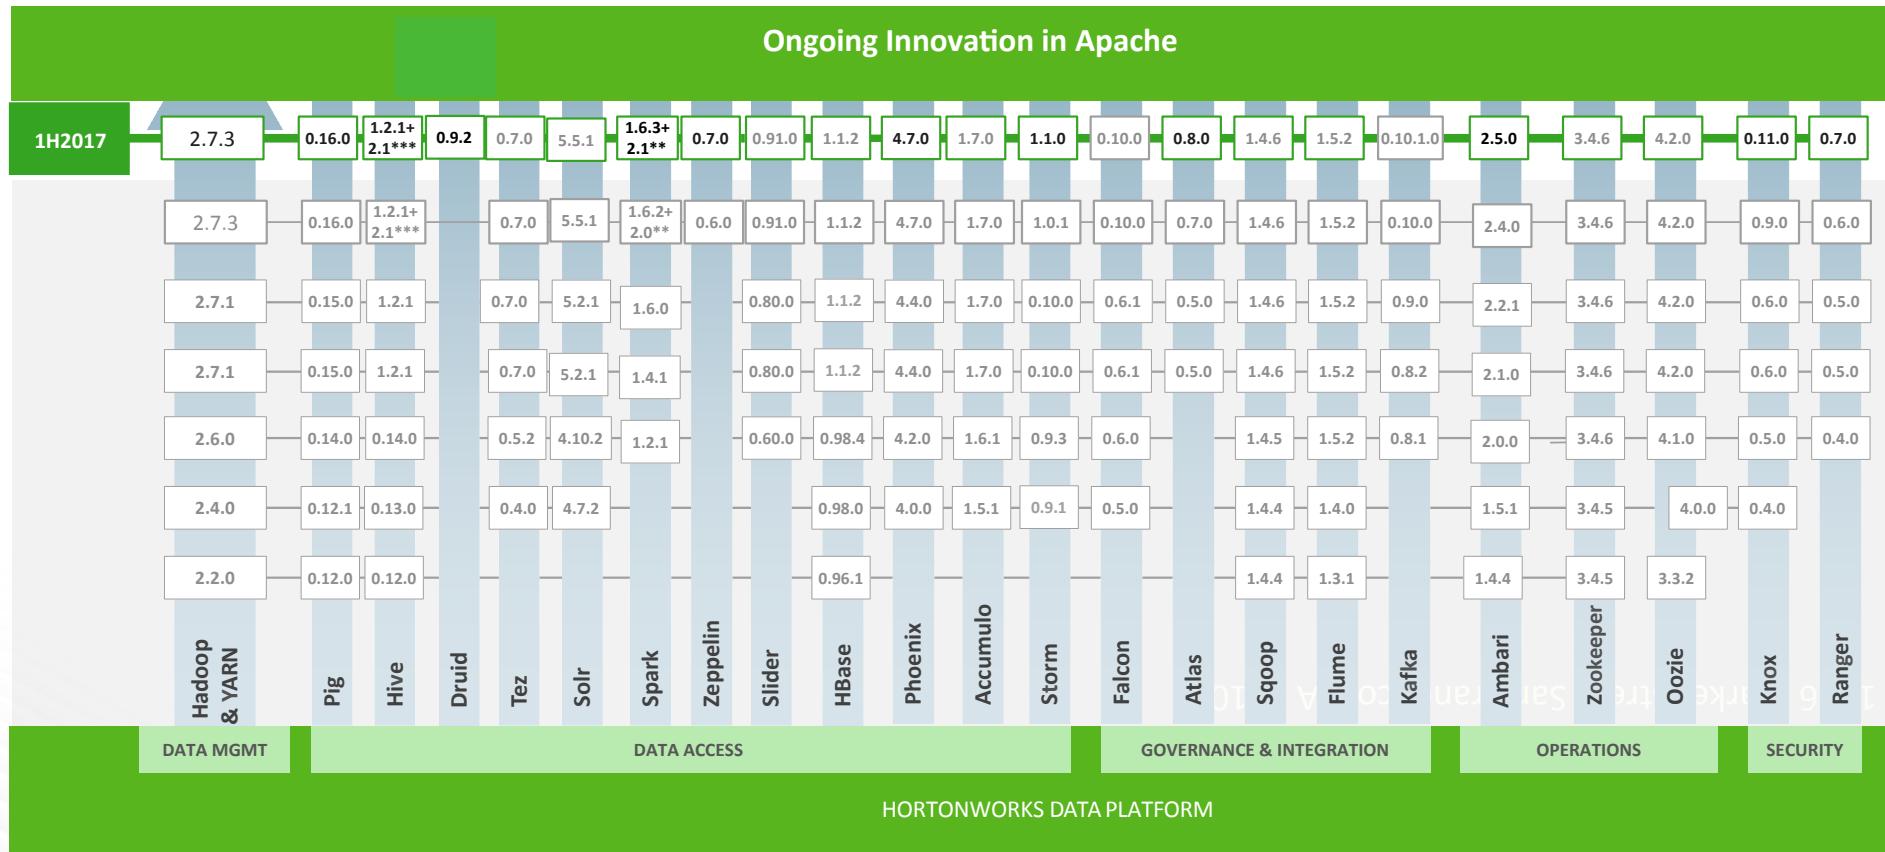
What is Apache Hadoop?

■ A brief history



What is Apache Hadoop?

■ A brief history: Versions



* HDP 2.6 – Shows current Apache branches being used. Final component version subject to change based on Apache release process.

** Spark 1.6.3+ Spark 2.1 – HDP 2.6 supports both Spark 1.6.3 and Spark 2.1 as GA

*** Hive 2.1 is GA within HDP 2.6.

What is Apache Hadoop?

- Open Source Apache Project
- Hadoop Core includes:
 - MapReduce: the processing part of Hadoop
 - Hadoop File System (HDFS): the data part of Hadoop
- Written in Java
- Runs on
 - Linux, Mac OS/X, Windows, and Solaris
 - Commodity hardware
- Key attributes
 - Redundant and reliable (no data loss)
 - Extremely powerful/scalable
 - Batch processing centric
 - Easy to program distributed applications

MapReduce 

HDFS

Machine

Hadoop Structure

■ What's running on the machine?

- The MapReduce server on a typical machine is called a TaskTracker.
- The HDFS server on a typical machine is called a DataNode.



TaskTracker

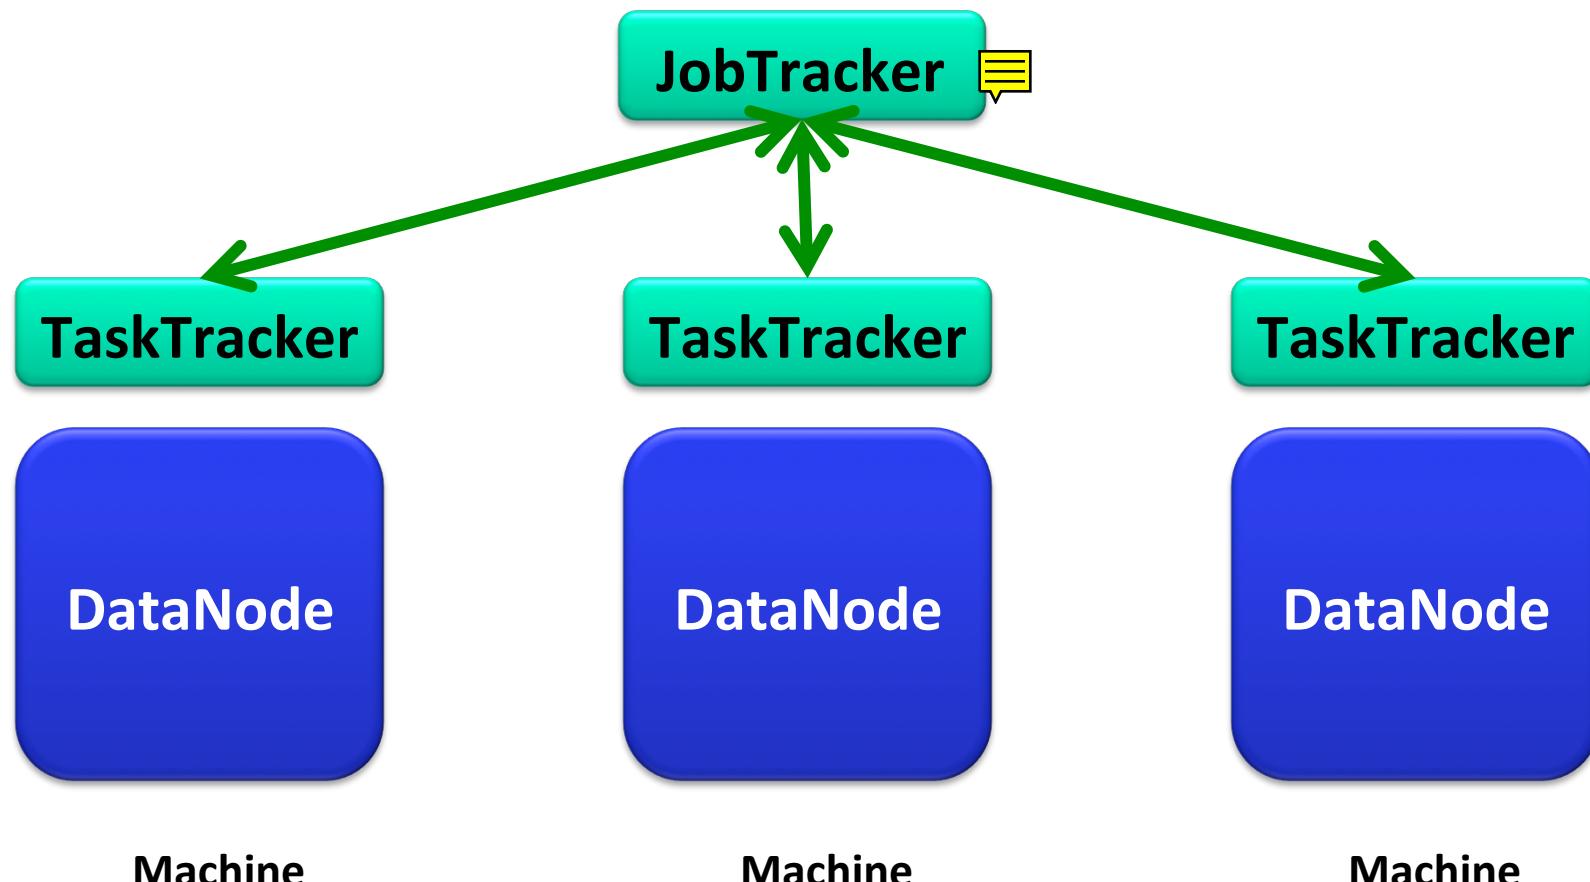
DataNode

Machine

Hadoop Structure

■ Hadoop cluster: MapReduce

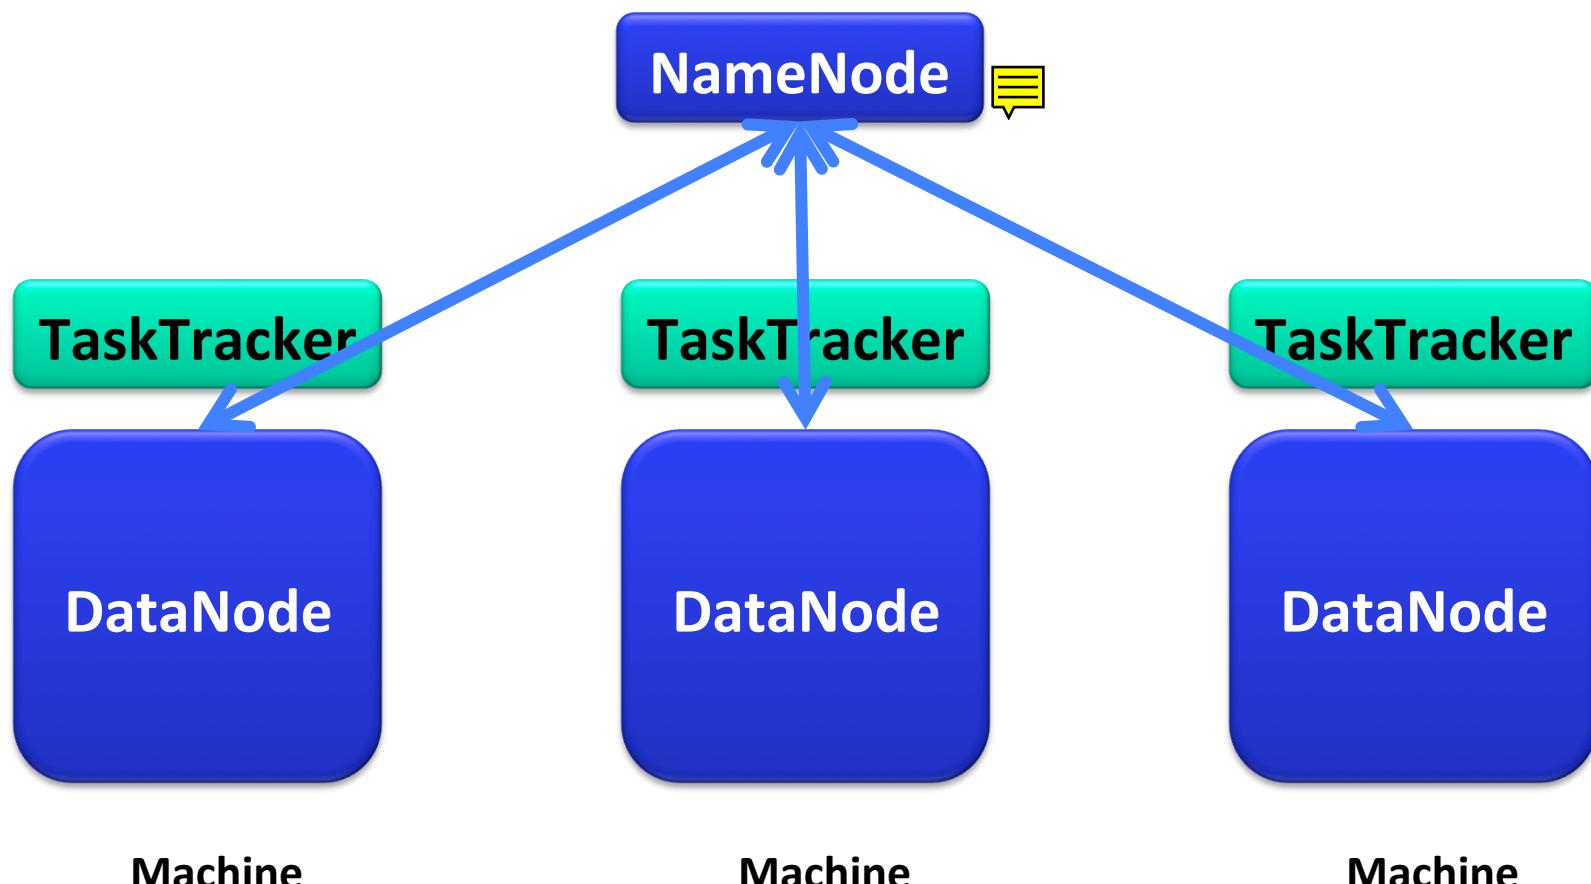
- Having multiple machines with Hadoop creates a cluster.
- JobTracker keeps track of jobs being run.



Hadoop Structure

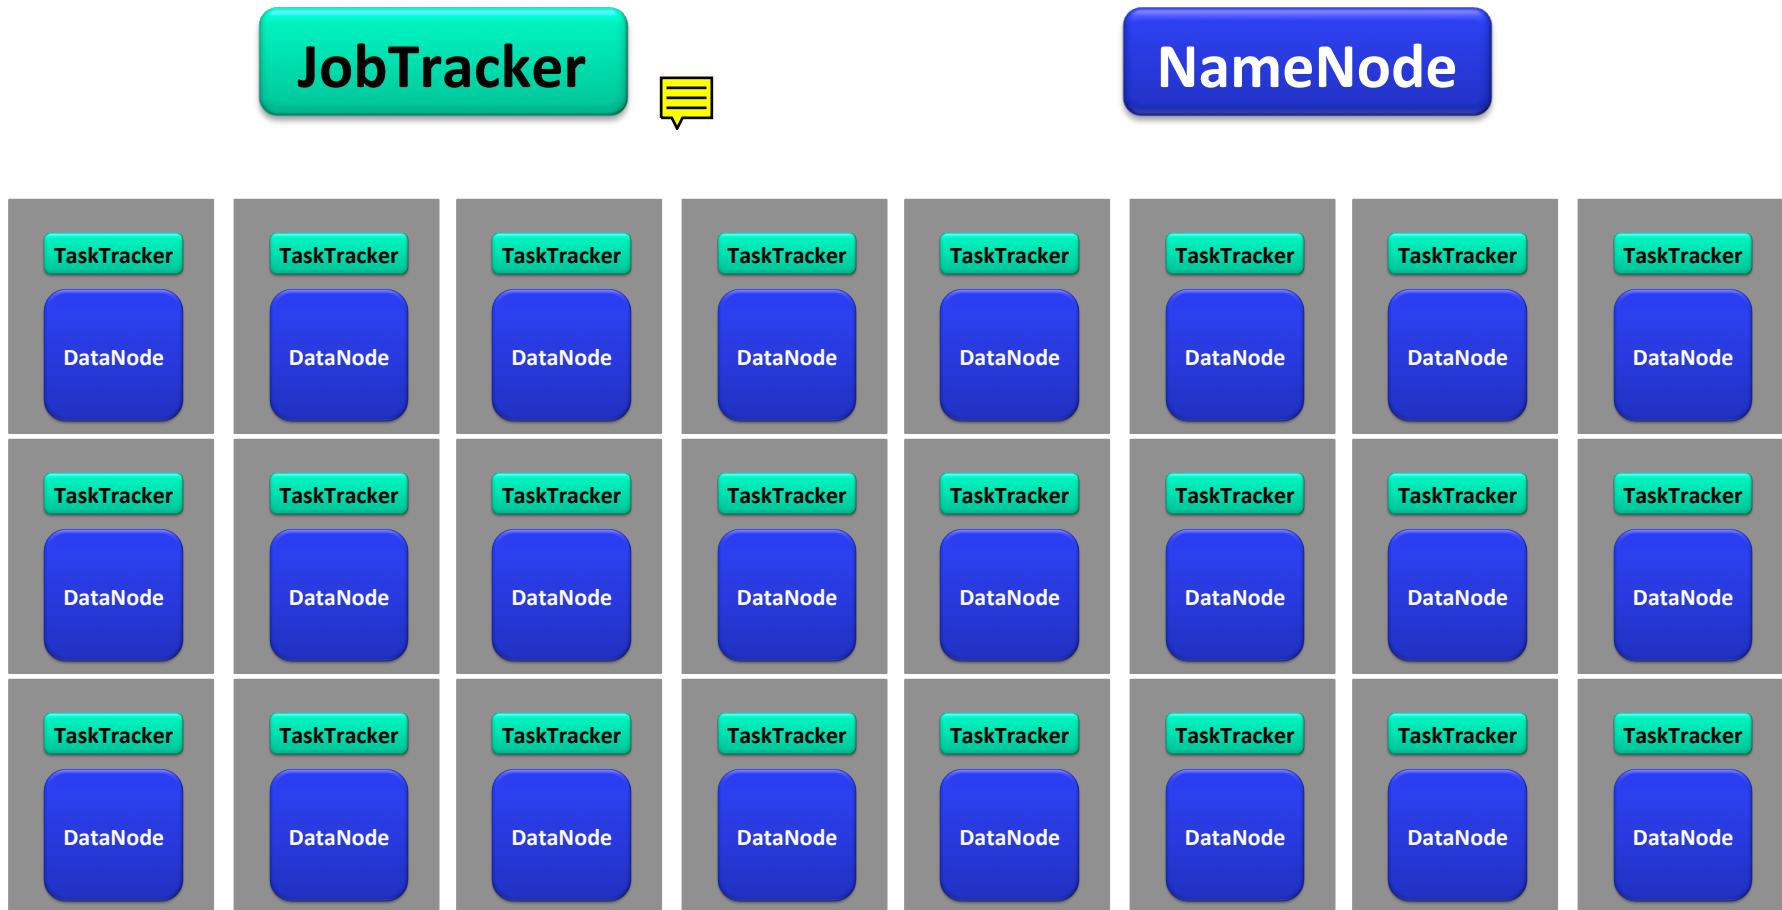
■ Hadoop cluster

- NameNode keeps information on data location



Hadoop Structure

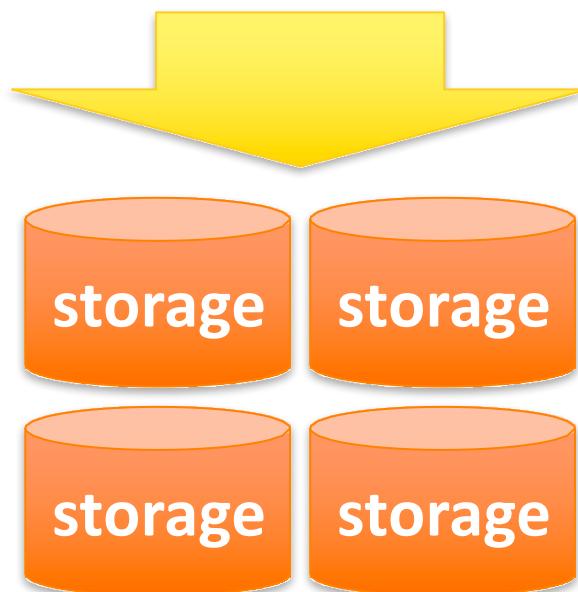
- Having many nodes..



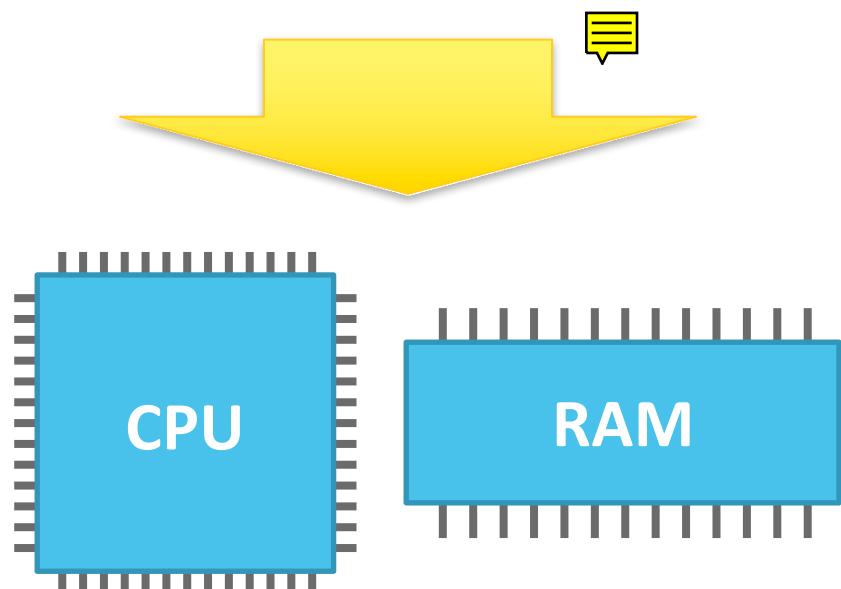
Hadoop Structure

- Apache Hadoop = Storage + Compute

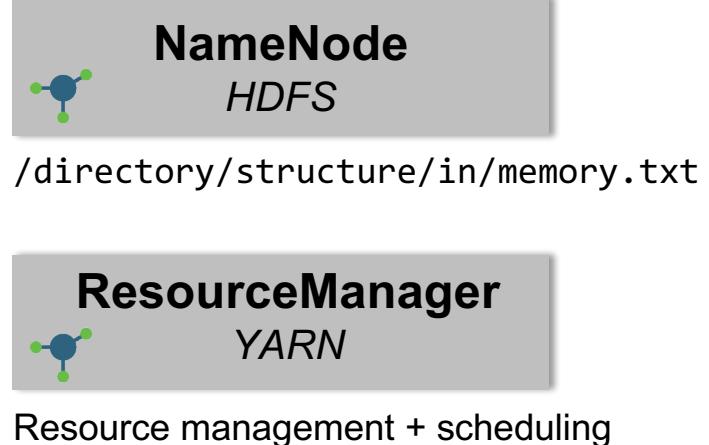
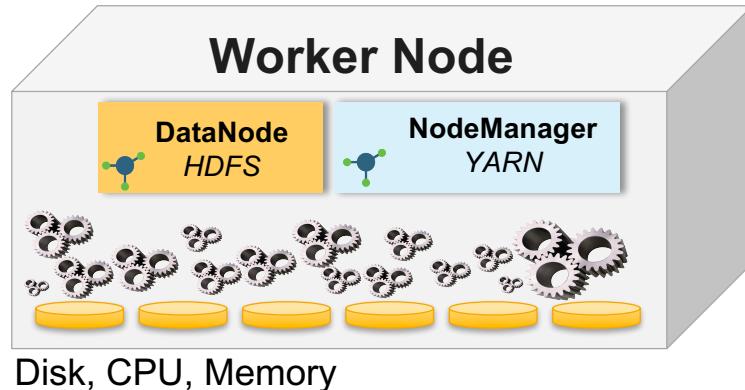
Hadoop Distributed File System (HDFS)



Yet Another Resource Negotiator (YARN)

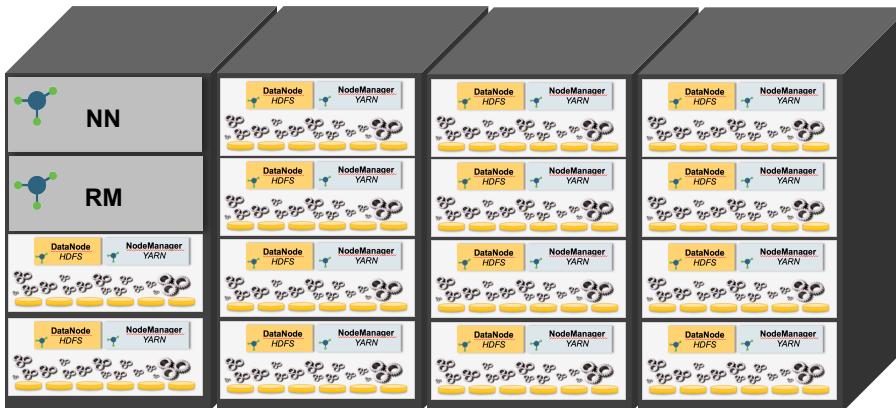


Hadoop is ...



Hadoop daemon

User application



Hadoop is ...



■ Reliable

- Data is typically held on multiple DataNodes
- Tasks that fail are redone

■ Scalable

- Same program runs on 1, 1000, or 4000 machines
- Scales linearly

■ Simple

- Having simple APIs

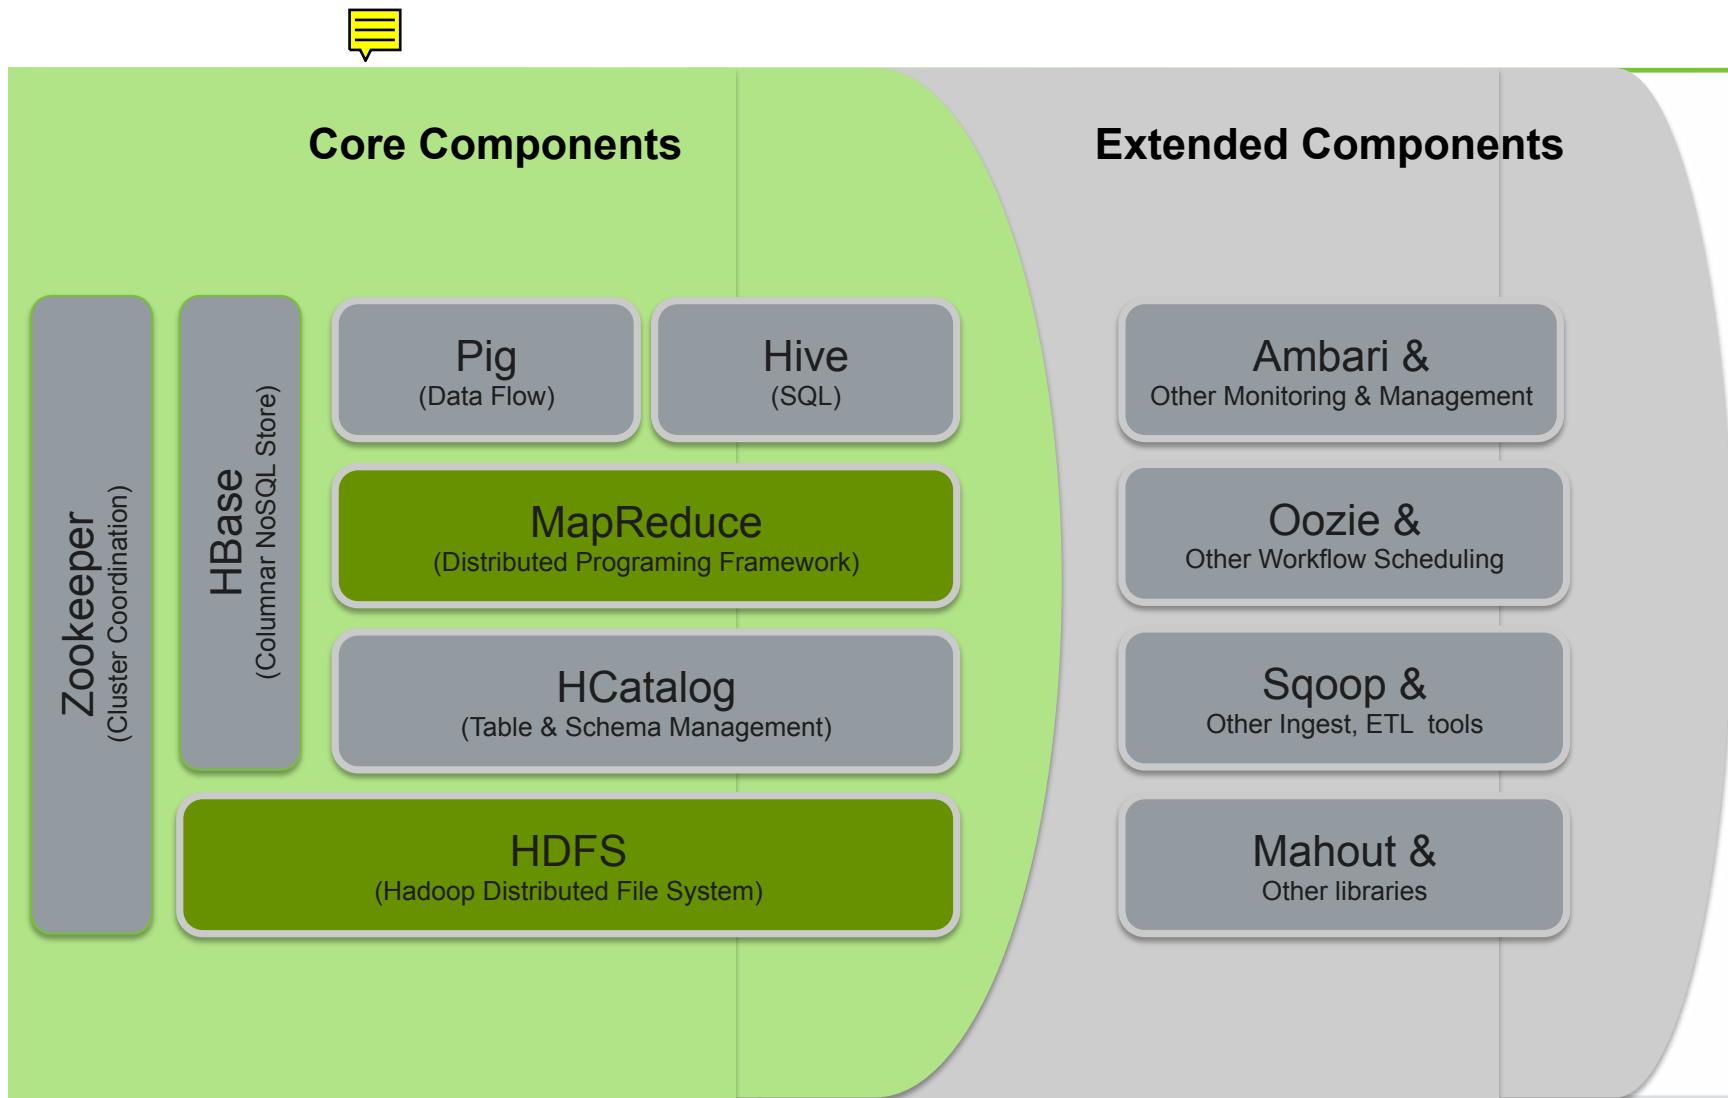
■ Very powerful

- You can process in parallel massive amounts of data
 - Petabytes of data
- Processing in parallel allows for the timely processing of massive amounts of data

Outline

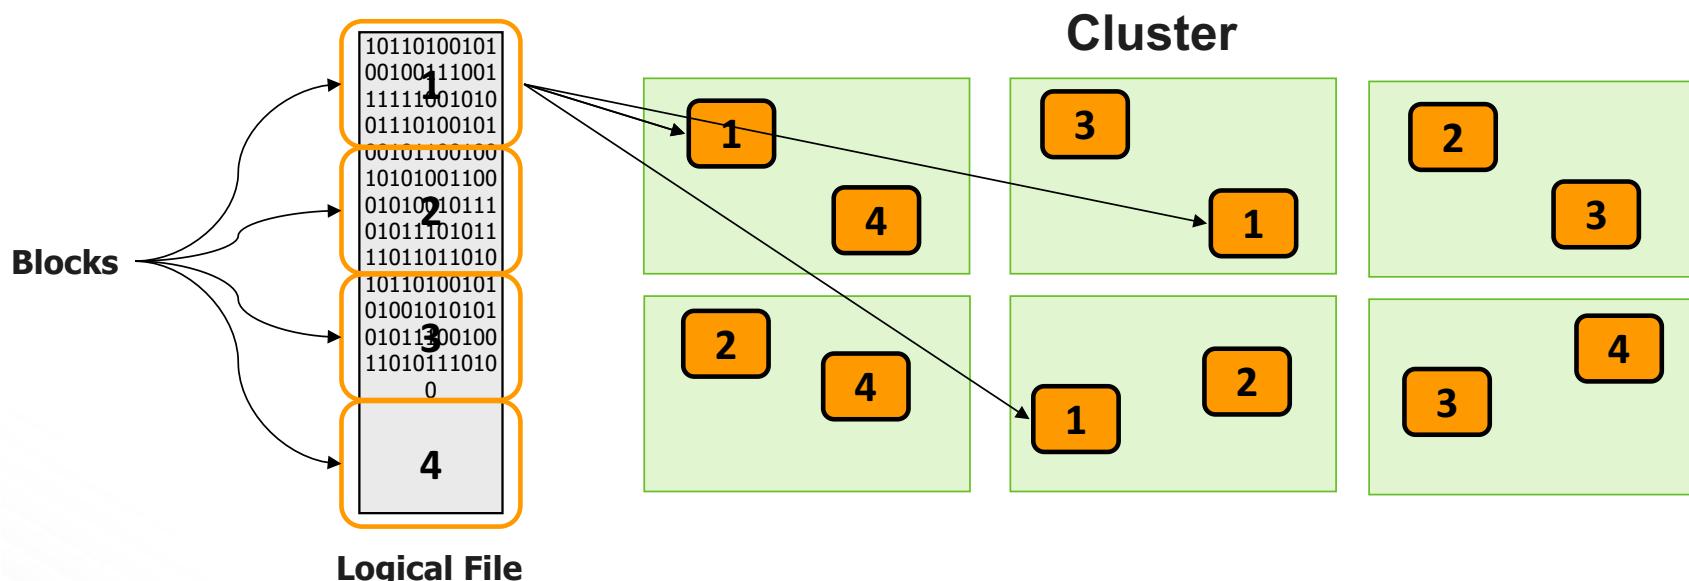
- What is Apache Hadoop?
- Key Components of Hadoop Stack
- Hadoop Distributed File System (HDFS)
- HDFS File System Shell Commands
- Apache Spark on HDFS

Key Components of Hadoop Stack



Key Components of Hadoop Stack

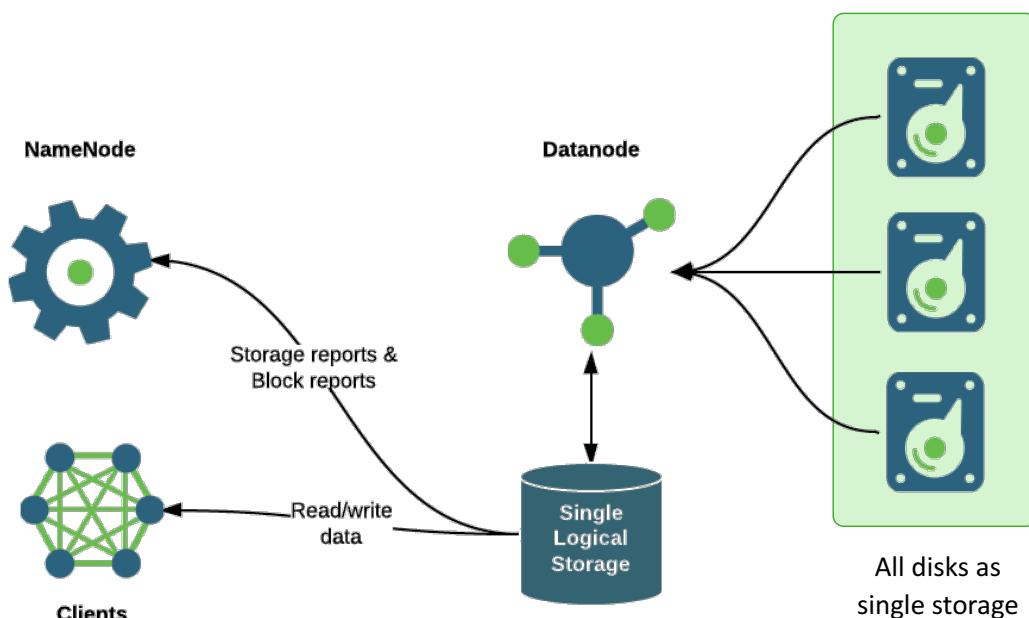
- **Hadoop Distributed File System (HDFS): Fault Tolerant Distributed Storage**
 - Divide files into big blocks and distribute 3 copies randomly across the cluster
 - Processing Data Locality
 - Not Just storage but computation



Key Components of Hadoop Stack

■ Hadoop Distributed File System (HDFS): Storage Architecture – Before

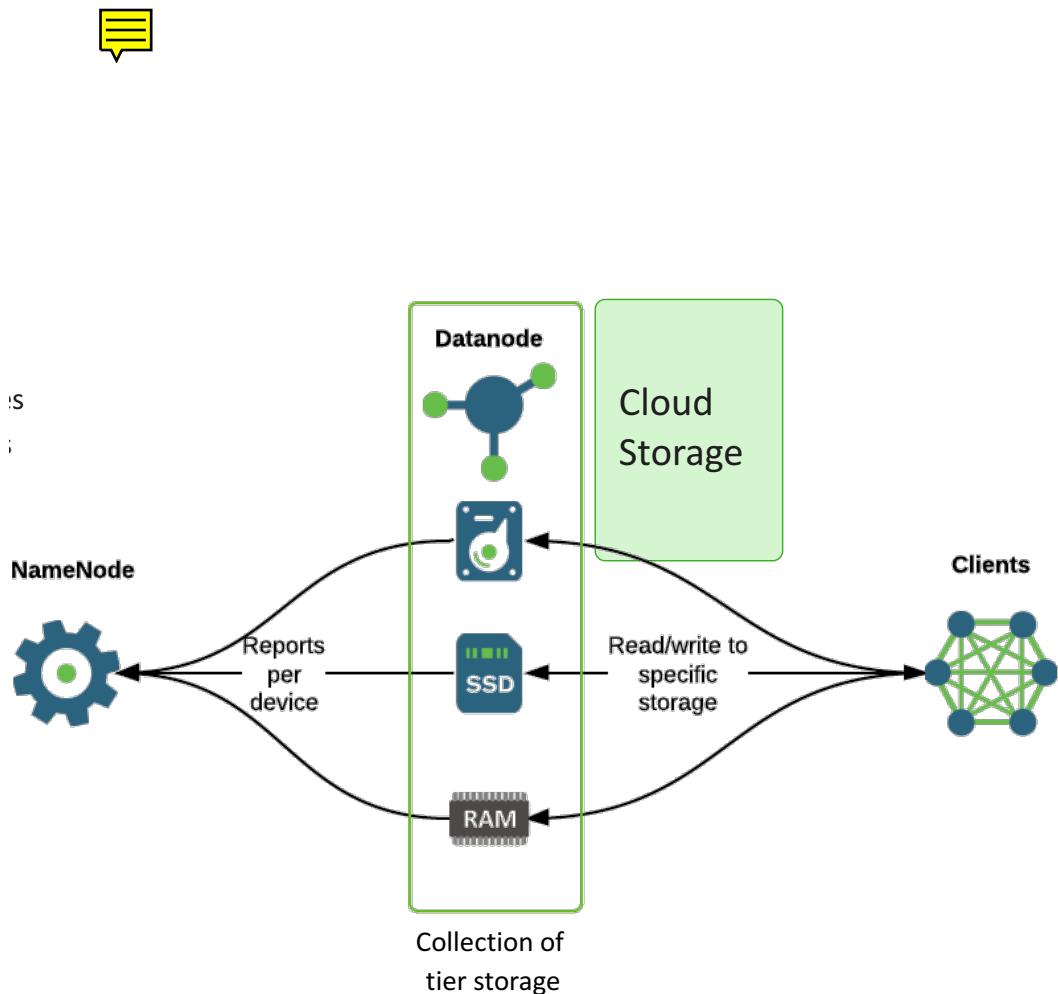
- DataNode is a single storage
- Storage is uniform - Only storage type Disk
- Storage types hidden from the file system



Key Components of Hadoop Stack

■ Hadoop Distributed File System (HDFS): Storage Architecture – Now

- New Architecture
 - DataNode is a collection of storages
 - Support different types of storages
 - Disk, SSDs, Memory
- Block Storage Policies
 - Describes how to store data blocks in HDFS



Key Components of Hadoop Stack

■ Hadoop Distributed File System (HDFS)

- It looks like a file system.

The screenshot shows a web-based interface for managing HDFS files. At the top, there's a header with a back button, a search bar, and a user profile icon. Below the header, the URL is shown as /user/it1/geolocation. To the right of the URL are buttons for 'New directory', 'Browse...', and a file upload area. A search bar labeled 'Search File Names' with a magnifying glass icon is also present.

The main area displays a table of files in the 'geolocation' directory. The columns are: Name, Size, Last Modified, Owner, Group, and Permission. The table includes two rows:

Name	Size	Last Modified	Owner	Group	Permission
geolocation.csv	514.3 kB	2016-03-13 19:42	maria_dev	hdfs	-rw-r--r--
trucks.csv	59.9 kB	2016-03-13 19:41	maria_dev	hdfs	-rw-r--r--

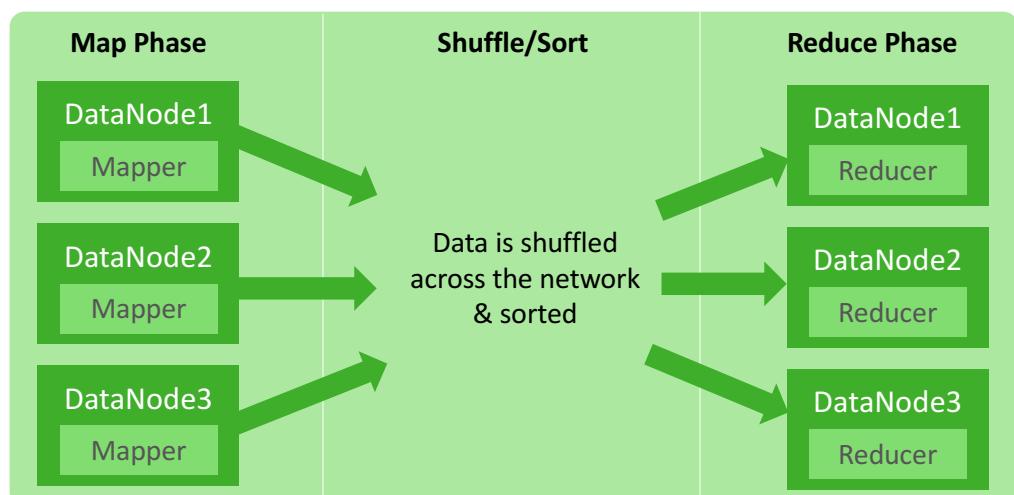
Below the table is a terminal window showing the output of the command 'hdfs dfs -ls /user/it1/geolocation'. The terminal output matches the file listing in the browser.

```
[it1@sandbox ~]$ hdfs dfs -ls /user/it1/geolocation
Found 2 items
-rw-r--r--  3 maria_dev hdfs      526677 2016-03-13 23:42 /user/it1/geolocation/geolocation.csv
-rw-r--r--  3 maria_dev hdfs      61378 2016-03-13 23:41 /user/it1/geolocation/trucks.csv
[it1@sandbox ~]$
```

Key Components of Hadoop Stack

■ MapReduce: Batch access to data

- Framework
 - Made for developing distributed applications to process vast amounts of data in-parallel on large clusters
- Proven
 - Reliable interface to Hadoop which works from GB to PB. But, batch oriented – Speed is not it's strong point.
- Ecosystem
 - Ported to Hadoop 2 to run on YARN. Supports original investments in Hadoop by customers and partner ecosystem.



Key Components of Hadoop Stack

■ MapReduce: What is it?

- Break a large problem into sub-solutions



Map

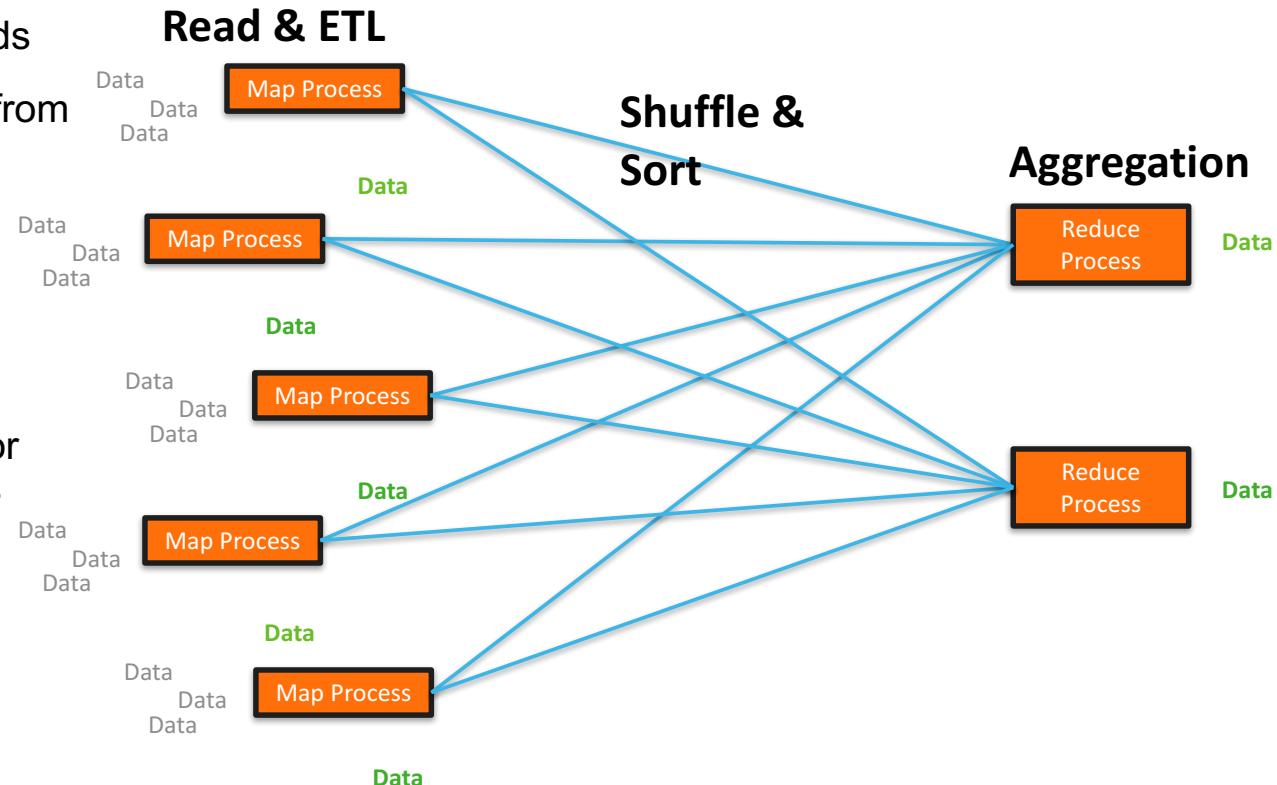
- Iterate over a large # of records
- Extract something of interest from each record

Shuffle

- Sort Intermediate results

Reduce

- Aggregate, summarize, filter or transform intermediate results
- Generate final output



* ETL: Extract, Transform, Load

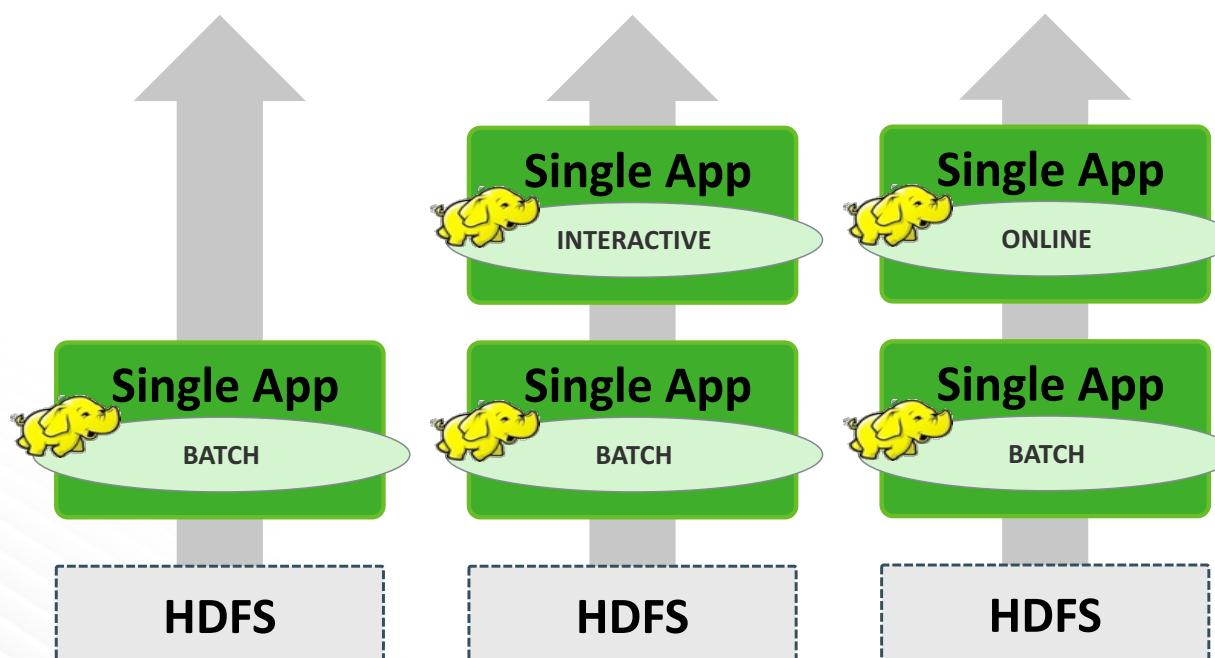
Key Components of Hadoop Stack

■ MapReduce: Key component for 1st Gen Hadoop

- Cost effective **batch** at scale



HADOOP 1.0 Built for Web-Scale Batch Apps



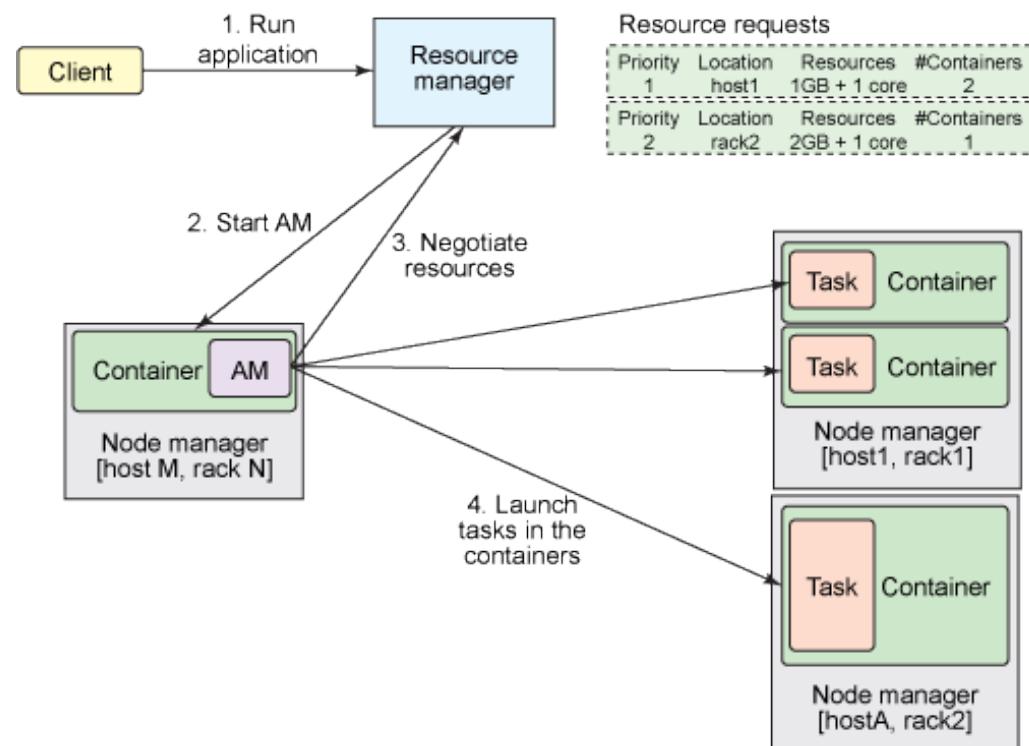
Key Components of Hadoop Stack

■ YARN: The resource manager for Apache Hadoop

■ Being able to share resources not only across MapReduce jobs, but also across the other processing frameworks

■ Correspondence (MapReduce -> YARN):

- Cluster manager -> ResourceManager
- JobTracker -> ApplicationMaster
- TaskTracker -> NodeManager
- MapReduce job -> Distributed application

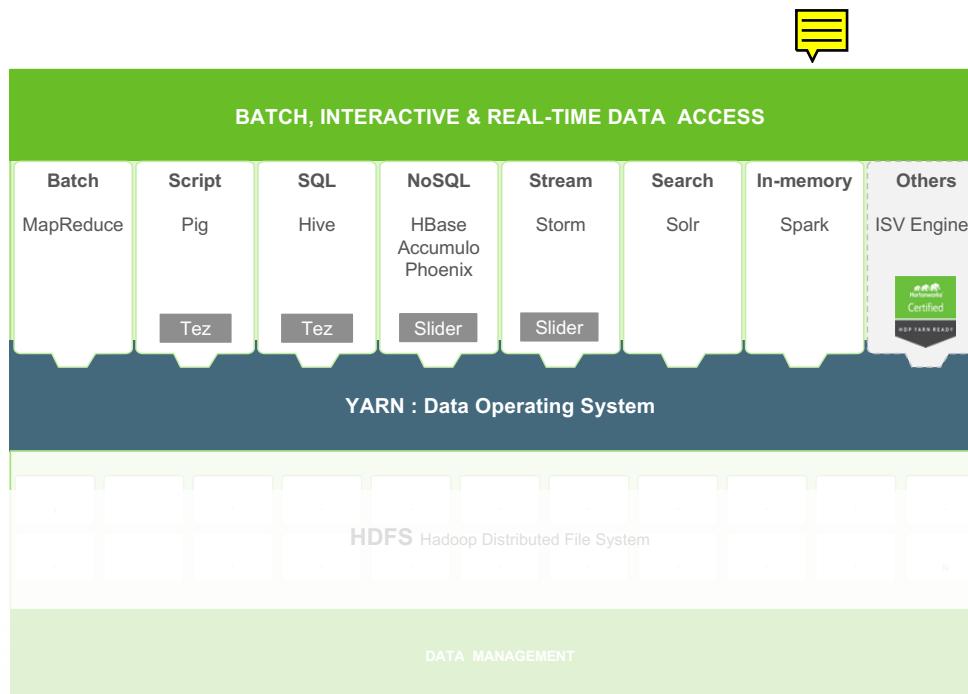


Source: IBM

Key Components of Hadoop Stack

■ YARN: Data Operating System

- Architectural center of Hadoop
 - Common data platform, many applications
 - Support multi-tenant access & processing
 - Batch, interactive & real-time use cases



Key Components of Hadoop Stack

■ YARN: Hadoop beyond batch!

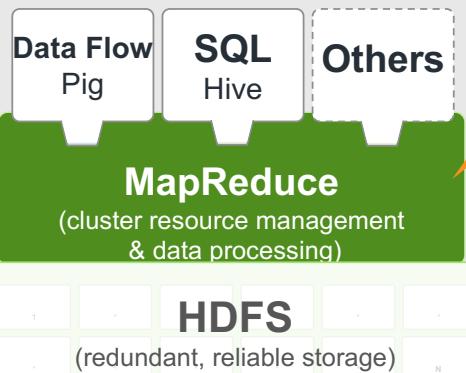
- A shift from the old to the new

Single Use System

Batch Apps

HADOOP 1

MapReduce as the Base

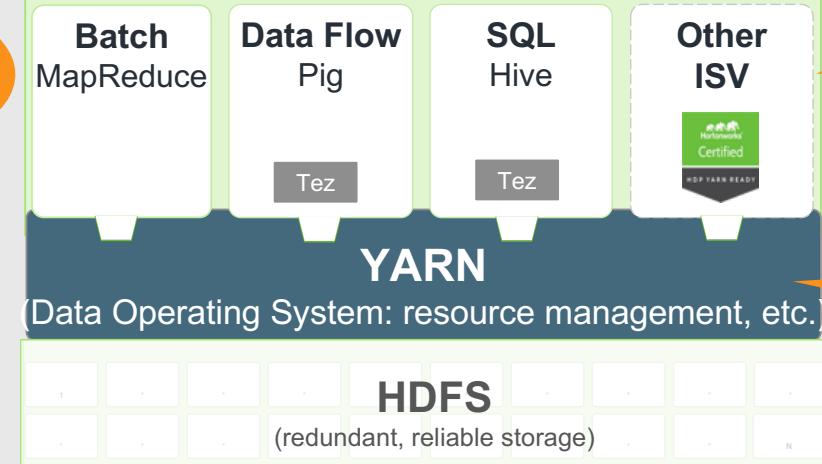


Multi Use Data Platform

Batch, Interactive, Online, Streaming, ...

HADOOP 2

Apache Yarn as a Base



Key Components of Hadoop Stack

■ Apache Hive: SQL in Hadoop



- Created by a team at Facebook
- Provides a standard SQL interface to data stored in Hadoop
- Quickly find value in raw data files
- Proven at petabyte scale
- Compatible with ALL major BI tools such as Tableau, Excel, MicroStrategy, Business Objects, etc...



Operational
/ MPP



Weblog



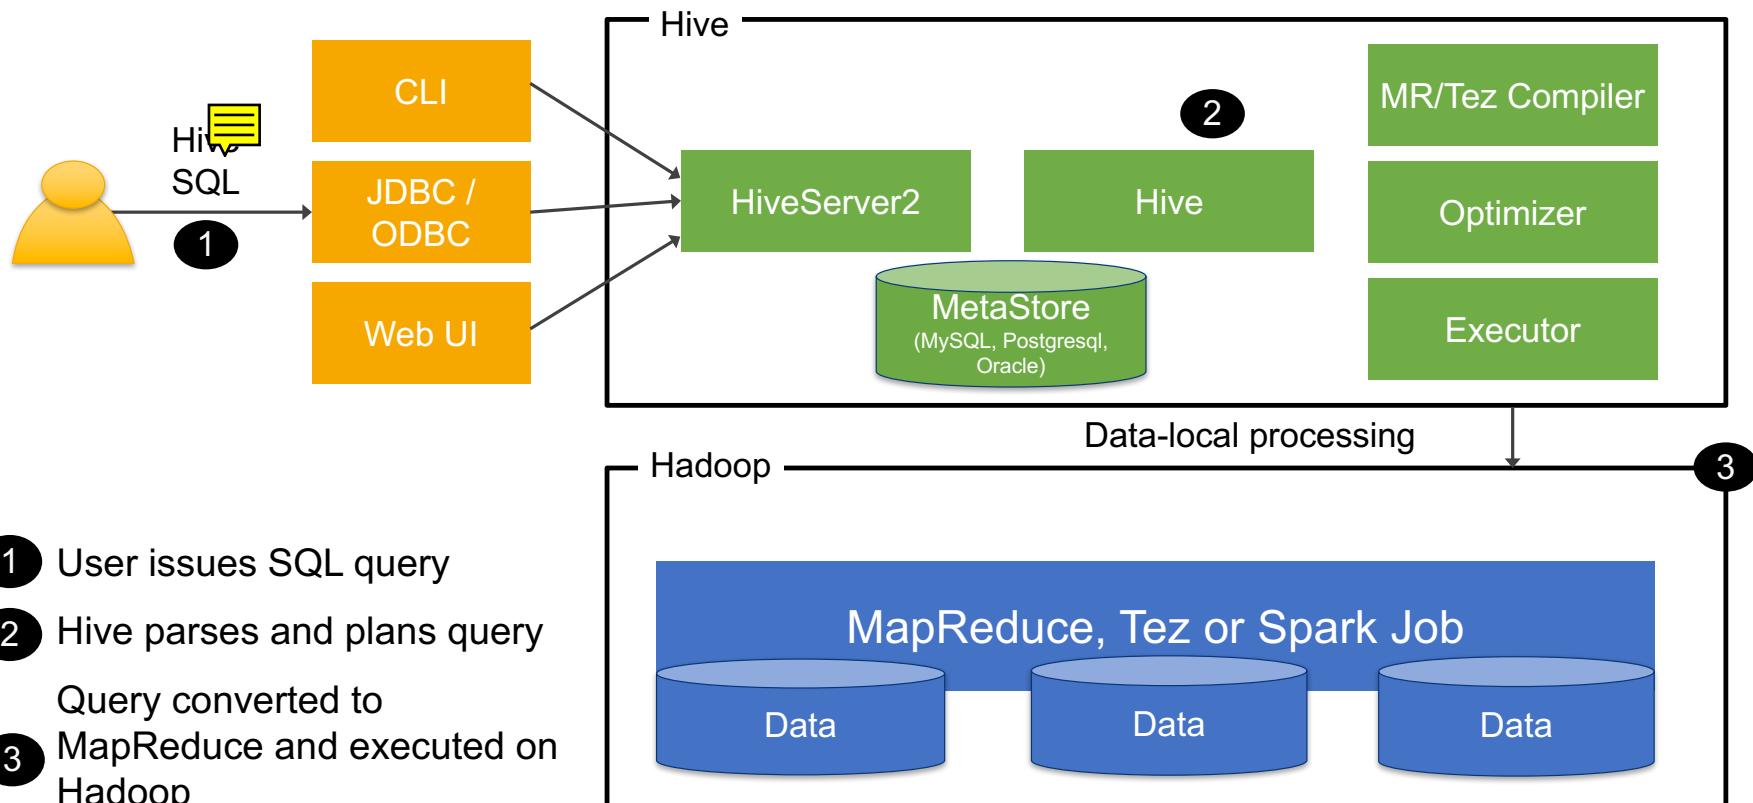
Mobile



Sensor

Key Components of Hadoop Stack

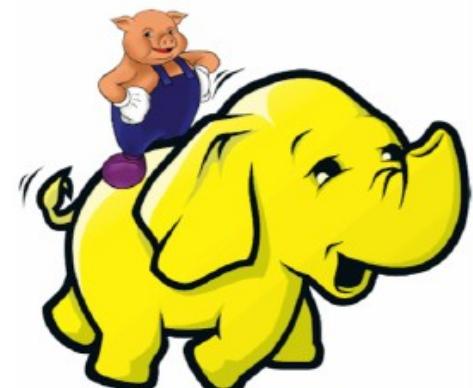
■ Apache Hive: Architecture



Key Components of Hadoop Stack

■ Apache Pig

- Data flow engine and scripting language (Pig Latin)
- Allows you to transform data and datasets
- Advantages over MapReduce
 - Reduces time to write jobs
 - Community support
 - Piggybank has a significant number of User-Defined Functions (UDF's) to help adoption
 - There are a large number of existing shops using PIG



Key Components of Hadoop Stack

■ Apache Pig: Example

- Maybe we want to join two datasets, from different sources, on a common value, and want to filter, and sort, and get top 5 sites

```
1 users = LOAD 'input/users' USING PigStorage(',')  
2   AS (name:chararray, age:int);  
3  
4 filtrd = FILTER users BY age >= 18 and age <= 25;  
5  
6 pages = LOAD 'input/pages' USING PigStorage(',')  
7   AS (user:chararray, url:chararray);  
8  
9 jnd = JOIN filtrd BY name, pages BY user;  
10  
11 grpdt = GROUP jnd BY url;  
12  
13 smmdt = FOREACH grpdt GENERATE group, COUNT(jnd) AS clicks;  
14  
15 srtdt = ORDER smmdt BY clicks DESC;  
16  
17 top5t = LIMIT srtdt 5;  
18  
19 STORE Top5t INTO 'output/top5sites' USING PigStorage(',');
```

Key Components of Hadoop Stack

■ HBase: Distributed column-oriented DB



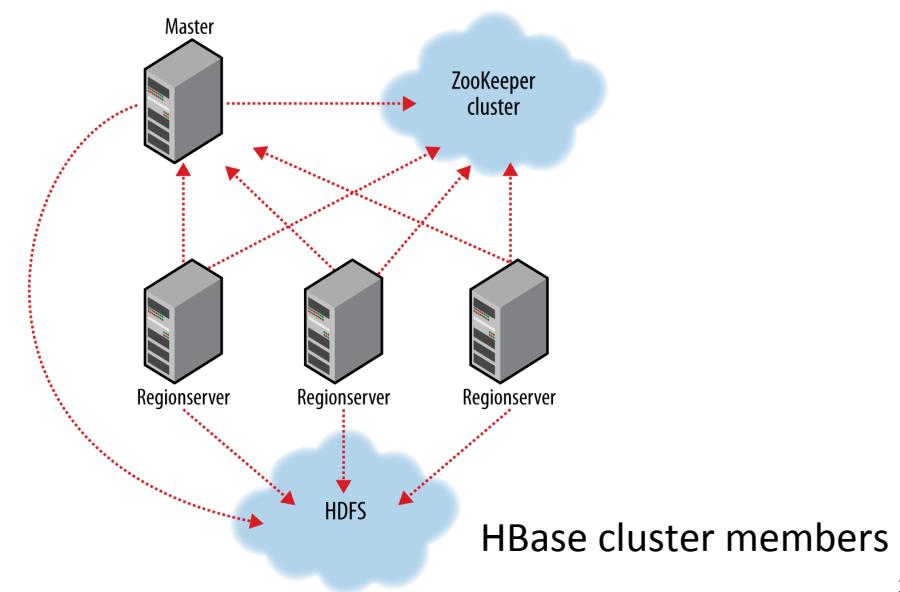
- Provides *real-time* read/write random access to very large datasets
 - Note that the original HDFS only provides append-only storage optimized for batch processing.
- Used for Facebook messages DB
- Unlike RDBMS it does not support SQL, but can host very large, sparsely populated tables well on commodity cluster

Increasing row key ↓

		Column family contents contents:image	Column family info info:format info:geo	
			jpeg	51.5,-0.1
			tiff	
000001				
000002			tiff	
000003			jpeg	51.8,-3.1

↑ Versions ↑ Cells

HBase data model (example: photo DB)

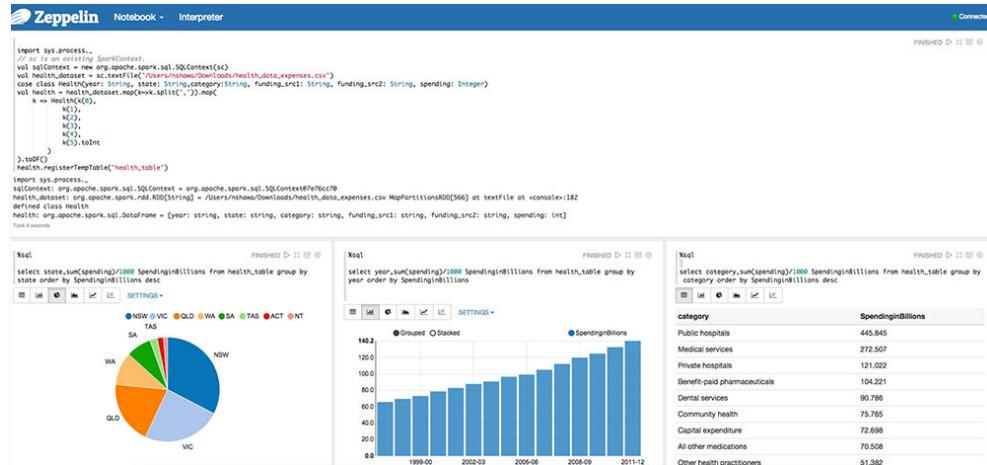


Key Components of Hadoop Stack

■ Apache Zeppelin



- Web-based notebook for data engineers, data analysts and data scientists
 - Brings interactive data ingestion, data exploration, visualization, sharing and collaboration features to Hadoop and Spark
- Modern data science studio
 - Scala with Spark
 - Python with Spark
 - SparkSQL
 - Apache Hive, and more.



Outline

- What is Apache Hadoop?
- Key Components of Hadoop Stack
- **Hadoop Distributed File System (HDFS)**
- HDFS File System Shell Commands
- Apache Spark on HDFS

HDFS

- Data is organized into files and directories
- Files are divided into uniform sized blocks (default 128MB) and distributed across cluster nodes
- HDFS exposes block placement so that computation can be migrated to data

HDFS

- **Blocks are replicated (default 3) to handle hardware failure**
- **Replication for performance and fault tolerance (Rack-Aware placement)**
- **HDFS keeps checksums of data for corruption detection and recovery**



HDFS

- Master-Worker Architecture
- Single NameNode
- Many (Thousands) DataNodes

HDFS Master (NameNode)



- Manages filesystem namespace
- File metadata (i.e. “inode”)
- Mapping inode to list of blocks + locations
- Authorization & Authentication
- Checkpoint & journal namespace changes

NameNode

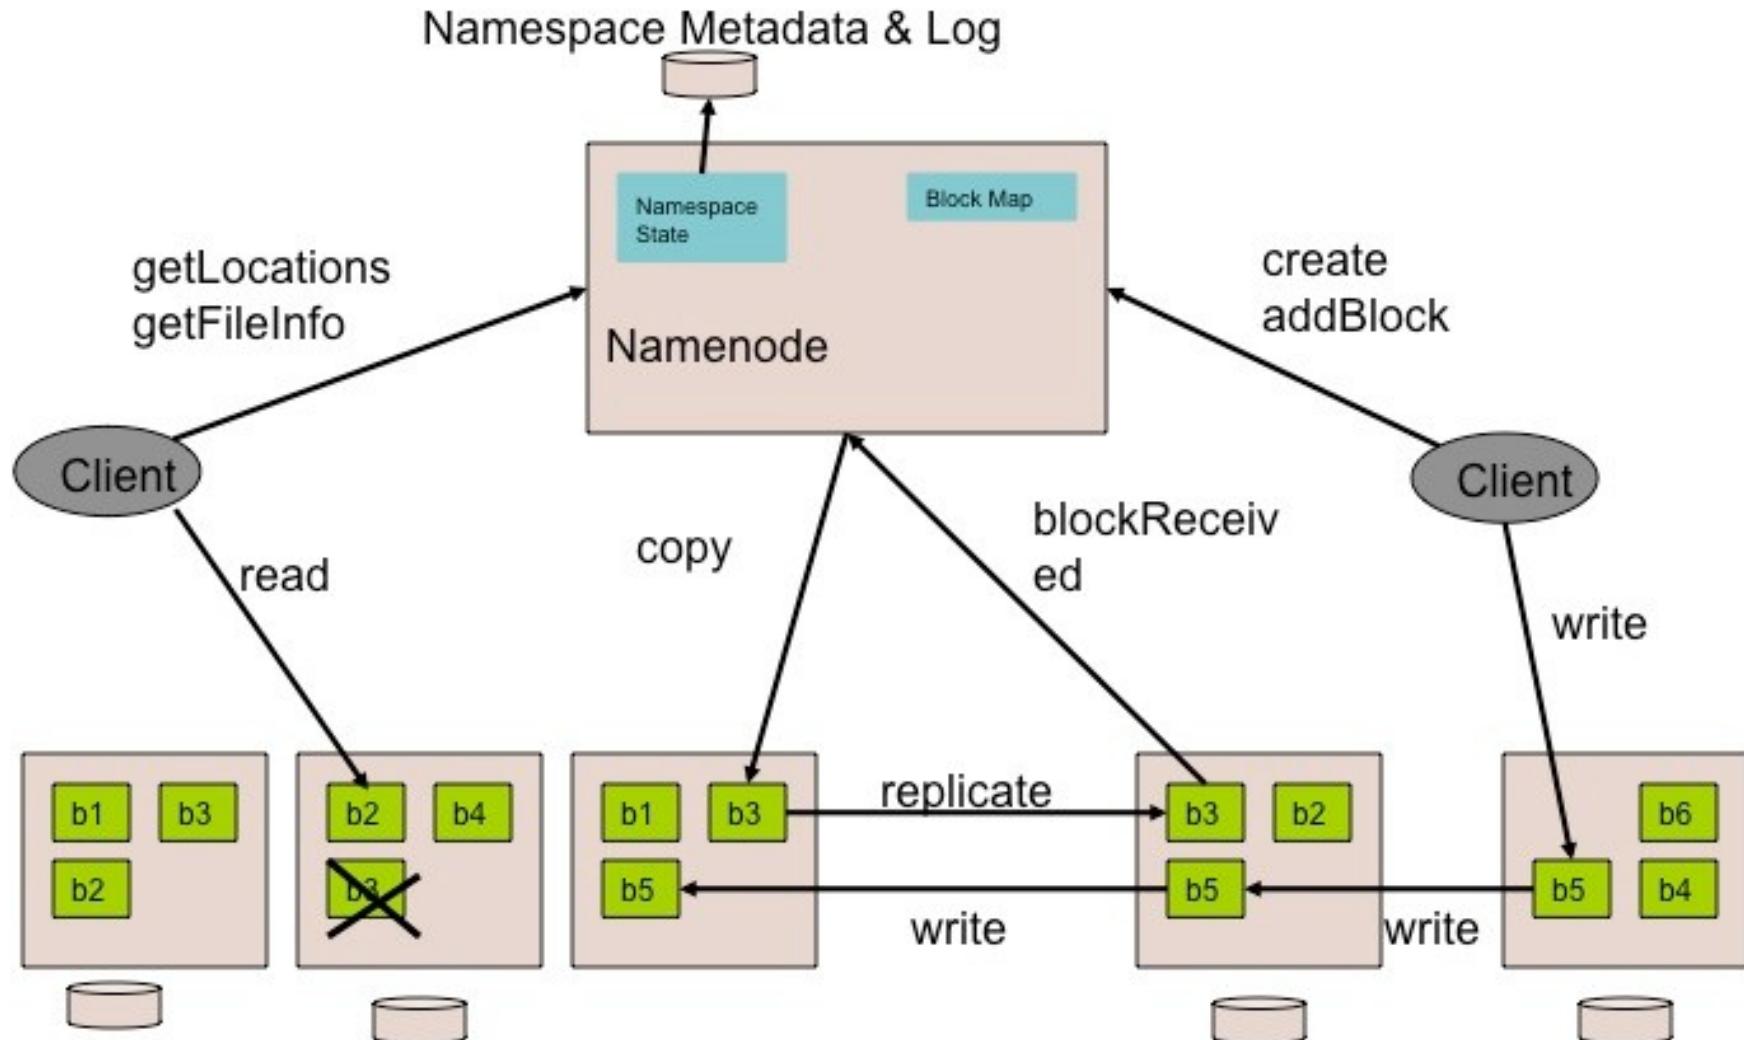


- Mapping of datanode to list of blocks
- Monitor datanode health
- Replicate missing blocks
- Keeps ALL namespace in memory
- 60M objects (File/Block) in 16GB

DataNodes

- Handle block storage on multiple volumes & block integrity
- Clients access the blocks directly from data nodes
- Periodically send heartbeats and block reports to NameNode
- Blocks are stored as underlying OS's files

HDFS Architecture



Design of HDFS

■ HDFS is ...



- "a filesystem designed for storing **very large files** with **streaming data access** patterns, running on clusters of **commodity hardware**"

■ Three key phrases in details

- Very large files
 - “Very large” in this context means files that are hundreds of megabytes, gigabytes, or terabytes in size

Design of HDFS

■ HDFS is ...

- "a filesystem designed for storing **very large files** with **streaming data access** patterns, running on clusters of **commodity hardware**"

■ Three key phrases in details

- Streaming data access
 - Built around the idea that the most efficient data processing pattern is a write-once, read-many-times pattern
 - A dataset is typically generated or copied from source, and then various analyses are performed on that dataset over time
 - Each analysis will involve a large proportion, if not all, of the dataset
 - the time to read the whole dataset is more important than the latency in reading the first record

Design of HDFS

■ HDFS is ...

- "a filesystem designed for storing **very large files** with **streaming data access** patterns, running on clusters of **commodity hardware**"

■ Three key phrases in details

- Commodity hardware
 - doesn't require expensive, highly reliable hardware
 - designed to run on clusters of commodity hardware (commonly available hardware that can be obtained from multiple vendors)
 - the chance of node failure across the cluster is high, at least for large clusters
 - designed to carry on working without a noticeable interruption to the user in the face of such failure

When HDFS Does NOT Work So Well?



■ For the following applications HDFS is NOT a good fit!

- Low-latency data access
 - Applications that require low-latency access to data, in the tens of milliseconds range, will not work well with HDFS (HDFS for high throughput, but not for low latency)
 - HBase is a better choice for the application
- Lots of small files
 - Namenode holds filesystem metadata in memory, the limit to the number of files in a filesystem is governed by the amount of memory on the namenode
 - As a rule of thumb, each file, directory, and block takes about 150 bytes
 - if you had one million files, each taking one block, you would need at least 300 MB of memory!

When HDFS Does NOT Work So Well?

- For the following applications HDFS is NOT a good fit! 
- Multiple writers, arbitrary file modifications
 - Files in HDFS may be written to by a single writer
 - Writes are always made at the end of the file, in append-only fashion
 - No support for multiple writers or for modifications at arbitrary offsets in the file

Example: Unigrams

- **Input: Huge text corpus**
 - Wikipedia Articles (40GB uncompressed)
- **Output: List of words sorted in descending order of frequency**

Unigrams

■ Implementation in Linux shell (if data fits in one node)

```
$ cat ~/wikipedia.txt | \  
sed -e 's/ /\n/g' | grep . | \  
sort | \  
uniq -c > \  
~/frequencies.txt
```



```
$ cat ~/frequencies.txt | \  
#  
cat | \  
sort -n -k1,1 -r | \  
#  
cat > \  
~/unigrams.txt
```

Unigrams

■ MapReduce for Unigrams: Pseudo-code (Phase #1)

- Counting # of occurrences for each word

```
mapper (filename, file-contents) :  
    for each word in file-  
    contents:  
        emit (word, 1)
```



```
reducer (word, values) :  
    sum = 0  
    for each value in values:  
        sum = sum + value  
    emit (word, sum)
```

Unigrams



■ MapReduce for Unigrams: Pseudo-code (Phase #2)

- Generating a sorted list

```
mapper (word, frequency) :  
    emit (frequency, word)  
  
reducer (frequency, words) :  
    for each word in words:  
        emit (word, frequency)
```

Unigrams

■ Java Mapper

- Just for reference (don't need to understand)

```
public static class MapClass extends MapReduceBase
    implements Mapper<LongWritable, Text, Text, IntWritable> {

    public void map(LongWritable key, Text value,
                    OutputCollector<Text, IntWritable> output,
                    Reporter reporter) throws IOException {

        String line = value.toString();
        StringTokenizer itr = new StringTokenizer(line);
        while (itr.hasMoreTokens()) {
            Text word = new Text(itr.nextToken());
            output.collect(word, new IntWritable(1));
        }
    }
}
```

Unigrams

■ Java Reducer

- Just for reference (don't need to understand)

```
public static class Reduce extends MapReduceBase
    implements Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> values,
                      OutputCollector<Text, IntWritable> output,
                      Reporter reporter) throws IOException {

        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

Unigrams

■ Driver (like main function)

```
public void run(String inputPath, String outputPath) throws
Exception
{
    JobConf conf = new JobConf(WordCount.class);
    conf.setJobName("wordcount");
    conf.setMapperClass(MapClass.class);
    conf.setReducerClass(Reduce.class);
    FileInputFormat.addInputPath(conf, new Path(inputPath));
    FileOutputFormat.setOutputPath(conf, new Path(outputPath));
    JobClient.runJob(conf);
}
```

Configuration

- **Unified Mechanism for**
 - Configuring Daemons
 - Runtime environment for Jobs/Tasks
- **Defaults: *-default.xml**
- **Site-Specific: *-site.xml**
- ***final* parameters**

Configuration

■ Example

```
<configuration>
    <property>
        <name>mapred.job.tracker</name>
        <value>head.server.node.com:9001</value>
    </property>
    <property>
        <name>fs.default.name</name>
        <value>hdfs://head.server.node.com:9000</value>
    </property>
    <property>
        <name>mapred.child.java.opts</name>
        <value>-Xmx512m</value>
        <final>true</final>
    </property>
    ....
</configuration>
```

Outline

- What is Apache Hadoop?
- Key Components of Hadoop Stack
- Hadoop Distributed File System (HDFS)
- **HDFS File System Shell Commands**
- Apache Spark on HDFS

HDFS File System Shell

■ HDFS File System (FS) shell

- includes various shell-like commands that directly interact with HDFS
- also support other file systems like local FS, HFTP FS, S3 FS, and others.
- is invoked by the following command:

```
$ bin/hadoop fs ...
```

- If HDFS is being used, the following command does the same.

```
$ bin/hdfs dfs ...
```

- If you need more information about supported shell commands:
 - Type **\$ bin/hdfs dfs -help**
 - Look up HDFS User Guide: <https://hadoop.apache.org/docs/stable/hadoop-project-dist/hadoop-common/FileSystemShell.html>

HDFS File System Shell Commands

■ ls: list files



- Usage: `hadoop fs -ls [-d] [-h] [-R] [<path> ...]`
- Desc: List the contents that match the specified file pattern. If path is not specified, the contents of /user/<currentUser> will be listed. -R recursively lists subdirectories.
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -ls /user/hadoop
Found 1 items
-rw-r--r--    1 ubuntu supergroup          14 2017-10-17 16:40 /
user/hadoop/file1
```



HDFS File System Shell Commands

■ cat: display file content

- Usage: `hadoop fs -cat [-ignoreCrc] <src> ...`
- Desc: Fetch all files that match the file pattern `<src>` and display their content on stdout
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -cat /user/hadoop/*
Hello world!
```

HDFS File System Shell Commands

■ mkdir: create a directory

- Usage: `hadoop fs -mkdir [-p] <path> ...`
- Desc: Create a directory in specified location. (-p Do not fail if the directory already exists).
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -mkdir /user/hadoop/mydir
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r--    1 ubuntu supergroup          14 2017-10-17 16:40 /
user/hadoop/file1
drwxr-xr-x    - ubuntu supergroup          0 2017-10-17 16:54 /
user/hadoop/mydir
```

HDFS File System Shell Commands

■ rm: delete files



- Usage: `hadoop fs -rm [-f] [-r|-R] [-skipTrash] <src> ...`
- Desc: Delete all files that match the specified file pattern. Equivalent to the Unix command "rm <src>". `-r` (or `-R`) recursively deletes directories.
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 16:40 /user/hadoop/file1
drwxr-xr-x - ubuntu supergroup          0 2017-10-17 16:54 /user/hadoop/mydir
$ bin/hadoop fs -rm -r /user/hadoop/mydir
17/10/17 16:59:11 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion
interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /user/hadoop/mydir
$ bin/hadoop fs -ls /user/hadoop
Found 1 items
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 16:40 /user/hadoop/file1
```

HDFS File System Shell Commands

■ put (or copyFromLocal): copy files from local file system

- Usage: hadoop fs -put [-f] [-p] [-l] <localsrc> ... <dst>
- Desc: Copy files from the local file system into fs. Copying fails if the file already exists, unless the -f flag is given.
- Example (assume we're in \$HADOOP_HOME directory)



```
# hello.txt and world.txt are in local file system
$ bin/hadoop fs -put hello.txt world.txt /user/hadoop
$ bin/hadoop fs -ls /user/hadoop
Found 3 items
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 16:40 /
user/hadoop/file1
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 17:09 /
user/hadoop/hello.txt
-rw-r--r-- 1 ubuntu supergroup          28 2017-10-17 17:09 /
user/hadoop/world.txt
```

HDFS File System Shell Commands

■ get (or copyToLocal): copy files from local file system

- Usage: `hadoop fs -get [-p] [-ignoreCrc] [-crc] <src> ... <localdst>`
- Desc: Copy files that match the file pattern `<src>` to the local name. `<src>` is kept. When copying multiple files, the destination must be a directory.
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -get /user/hadoop/hello.txt hello.copy.txt
$ md5sum hello.txt hello.copy.txt # compare two versions
5c06f4b1a10b425fb707d16bb4d5d9a9  hello.txt
5c06f4b1a10b425fb707d16bb4d5d9a9  hello.copy.txt
```

HDFS File System Shell Commands

■ **getmerge: get files and merge them into local file system**

- Usage: `hadoop fs -getmerge [-nl] <src> <localdst>`
- Desc: Get all the files in the directories that match the source file pattern and merge and sort them to only one file on local fs. `<src>` is kept. `-nl` adds a newline character at the end of each file
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -getmerge /user/hadoop/*.txt helloworld.txt  
$ cat helloworld.txt    # from local file system  
Hello world!
```

To the world: Hello there!

HDFS File System Shell Commands



■ **appendToFile: appends the contents of local files to dst file**

- Usage: `hadoop fs -appendToFile <localsrc> ... <dst>`
- Desc: Appends the contents of all the given local files to the given dst file. The dst file will be created if it does not exist. If <localSrc> is -, then the input is read from stdin.
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -appendToFile world.txt /user/hadoop/hello.txt  
$ bin/hadoop fs -cat hello.txt  
Hello world!
```

To the world: Hello there!

HDFS File System Shell Commands

■ mv: move files

- Usage: `hadoop fs -mv <src> ... <dst>`
- Desc: Move files that match the specified file pattern `<src>` to a destination `<dst>`. When moving multiple files, the destination must be a directory.
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -mv /user/hadoop/hello.txt /user/hadoop/olleh.txt
$ bin/hadoop fs -ls /user/hadoop
Found 3 items
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 16:40 /
user/hadoop/file1
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 17:09 /
user/hadoop/olleh.txt
-rw-r--r-- 1 ubuntu supergroup          28 2017-10-17 17:09 /
user/hadoop/world.txt
```

HDFS File System Shell Commands

■ chown: change ownership

- Usage: `hadoop fs -chown [-R] [OWNER][:[GROUP]] PATH...`
- Desc: Changes owner and group of a file. This is similar to the shell's `chown` command with a few exceptions. `-R` modifies the files recursively.
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r-- 1 ubuntu supergroup          14 2017-10-17 17:09 /
user/hadoop/hello.txt
-rw-r--r-- 1 ubuntu supergroup          28 2017-10-17 17:09 /
user/hadoop/world.txt
$ bin/hadoop fs -chown jaewooklee /user/hadoop/*
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r-- 1 jaewooklee supergroup      14 2017-10-17
17:09 /user/hadoop/hello.txt
-rw-r--r-- 1 jaewooklee supergroup      28 2017-10-17
17:09 /user/hadoop/world.txt
```

HDFS File System Shell Commands

■ chmod: change permissions



- Usage: `hadoop fs -chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...`
- Desc: Changes permissions of a file. This works similar to the shell's chmod command with a few exceptions. `-R` modifies the files recursively.
- Example (assume we're in `$HADOOP_HOME` directory)



```
$ bin/hadoop fs -chmod -R g+rwx /user/hadoop
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-rwrxr-- 1 jaewooklee supergroup          42 2017-10-18
00:19 /user/hadoop/hello.txt
-rw-rwrxr-- 1 jaewooklee supergroup          28 2017-10-17
17:09 /user/hadoop/world.txt
```

HDFS File System Shell Commands

■ **expunge: empty trash**



- Usage: `hadoop fs -expunge`
- Desc: Delete files from the trash that are older than the retention threshold
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -expunge
17/10/17 17:33:57 INFO fs.TrashPolicyDefault: Namenode trash
configuration: Deletion interval = 0 minutes, Emptier interval =
0 minutes.
$ # trash emptied
```

HDFS File System Shell Commands

■ du: display length of files

- Usage: hadoop fs -du [-s] [-h] <path> ...
- Desc: Show the amount of space, in bytes, used by the files that match the specified file pattern. The output is in the form: size – name (path)
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -du /user/hadoop
14  /user/hadoop/file1
14  /user/hadoop/olleh.txt
28  /user/hadoop/world.txt
```

HDFS File System Shell Commands

■ setrep: change replication factor of file

- Usage: `hadoop fs -setrep [-R] [-w] <rep> <path> ...`
- Desc: Set the replication level of a file. If `<path>` is a directory then the command recursively changes the replication factor of all files under the directory tree rooted at `<path>`. `-w` waits until the replication to complete (this potentially takes a very long time).
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -setrep -R 3 /user/hadoop
Replication 3 set: /user/hadoop/hello.txt
Replication 3 set: /user/hadoop/world.txt
$ bin/hadoop fs -ls /user/hadoop
Found 2 items
-rw-r--r--    3 ubuntu supergroup          42 2017-10-18 00:19 /
user/hadoop/hello.txt
-rw-r--r--    3 ubuntu supergroup         28 2017-10-17 17:09 /
user/hadoop/world.txt
```

HDFS File System Shell Commands

■ stat: show stat information

- Usage: hadoop fs -stat [format] <path> ...
- Desc: Print statistics about the file/directory at <path> in the specified format. Format accepts filesize in blocks (%b), type (%F), group name of owner (%g), name (%n), block size (%o), replication (%r), user name of owner (%u), modification date (%y, %Y).
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -stat %b::%F::%g::%n::%o::%r::%y /user/hadoop/  
world.txt  
28::regular file::supergroup::world.txt:134217728::3::2017-10-17  
17:09:56
```

HDFS File System Shell Commands

■ tail: display last part to stdout



- Usage: `hadoop fs -tail [-f] <file>`
- Desc: Show the last 1KB of the file. -f shows appended data as the file grows.
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -tail /user/hadoop/hello.txt  
Hello world!
```

```
To the world: Hello there!
```

HDFS File System Shell Commands

■ test: check attribute of file/dir

- Usage: `hadoop fs -test -[defsz] <path>`
- Desc: Answer various questions about `<path>`, with result via exit status.
(`-d`: return 0 if `<path>` is a directory, `-e`: return 0 if `<path>` exists, and so on.)
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -test -e /user/hadoop/hello.txt # returns 0  
$ bin/hadoop fs -test -e /user/hadoop/hi.txt      # returns 1
```

HDFS File System Shell Commands

■ text: outputs file in text format

- Usage: hadoop fs -text [-ignoreCrc] <src> ...
- Desc: Takes a source file and outputs the file in text format. The allowed formats are zip and TextRecordInputStream and Avro.
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -text /user/hadoop/hello.zip
P|RKY???"*      hello.txtUT      ???Y???Yux
                           ???H???W(??/?IQ??
?W(??H?p?<?r@??T?P|RKY???"*      ??hello.txtUT???Yux
                           ??PKOeu
```

HDFS File System Shell Commands

■ touchz: create a zero-length file

- Usage: `hadoop fs -touchz <path> ...`
- Desc: Creates a file of zero length at `<path>` with current time as the timestamp of that `<path>`. An error is returned if the file exists with non-zero length.
- Example (assume we're in `$HADOOP_HOME` directory)

```
$ bin/hadoop fs -touchz /user/hadoop/zerolengh.txt
$ bin/hadoop fs -ls /user/hadoop
Found 3 items
-rw-r--r--    3 ubuntu supergroup          42 2017-10-18 00:19 /
user/hadoop/hello.txt
-rw-r--r--    3 ubuntu supergroup          28 2017-10-17 17:09 /
user/hadoop/world.txt
-rw-r--r--    1 ubuntu supergroup          0  2017-10-18 00:56 /
user/hadoop/zerolengh.txt
```

HDFS File System Shell Commands

■ count: count files, dirs, and bytes

- Usage: hadoop fs -count [-q] [-h] <path> ...
- Desc: Count the number of directories, files and bytes under the paths that match the specified file pattern. The output columns are:
DIR_COUNT FILE_COUNT CONTENT_SIZE FILE_NAME or
(with -q flag) QUOTA REMAINING_QUOTA SPACE_QUOTA
REMAINING_SPACE_QUOTA DIR_COUNT FILE_COUNT CONTENT_SIZE
FILE_NAME
- Example (assume we're in \$HADOOP_HOME directory)

```
$ bin/hadoop fs -count -q /user/hadoop
    none          inf          none          inf
    1            3            70  /user/hadoop
```

HDFS Commands

■ fsck: File system check

- Runs the HDFS filesystem checking utility

```
% hdfs fsck /
Connecting to namenode via http://localhost:50070/fsck?ugi=ubuntu&path=%2F
FSCK started by ubuntu (auth:SIMPLE) from /127.0.0.1 for path / at Tue Oct 17
17:44:22 UTC 2017
...Status: HEALTHY
Total size:      56 B
Total dirs:      3
Total files:     3
Total symlinks:   0
Total blocks (validated): 3 (avg. block size 18 B)
Minimally replicated blocks: 3 (100.0 %)
Over-replicated blocks:    0 (0.0 %)
Under-replicated blocks:   0 (0.0 %)
Mis-replicated blocks:     0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks:        0
Missing replicas:       0 (0.0 %)
Number of data-nodes:    1
Number of racks:        1
FSCK ended at Tue Oct 17 17:44:22 UTC 2017 in 3 milliseconds

The filesystem under path '/' is HEALTHY
```

HDFS Commands

■ **distcp: distributed copy**



- Usage: hadoop distcp [options] <srcurl> <desturl>
- It is used for large inter/intra-cluster copying. It uses MapReduce to effect its distribution, error handling and recovery, and reporting. It expands a list of files and directories into input to map tasks, each of which will copy a partition of the files specified in the source list.
- Reference: [https://hadoop.apache.org/docs/stable/hadoop-distcp/ DistCp.html](https://hadoop.apache.org/docs/stable/hadoop-distcp/DistCp.html)

```
$ bin/hadoop distcp hdfs://nn1:8020/foo/bar hdfs://nn2:8020/bar/foo
```

Outline

- What is Apache Hadoop?
- Key Components of Hadoop Stack
- Hadoop Distributed File System (HDFS)
- HDFS File System Shell Commands
- Apache Spark on HDFS

Collocating Spark and HDFS

- **Spark works well with HDFS.**
 - Gives applications a chance to achieve data locality for HDFS running on the same machines
- **Using Spark with HDFS is simple:**
 - Just specifying `hdfs://master:port/path` for your input and output suffices.
- **Caveat – version mismatch**
 - Spark and HDFS (Hadoop) versions should match!
 - The HDFS protocol changes across Hadoop versions.
 - If you build Spark from source, you can specify the Hadoop version number (different from the default version) by setting `SPARK_HADOOP_VERSION` environment variable.

Accessing HDFS from Spark

■ Example: Count the total words in text files

```
1 >>> text_file = sc.textFile("hdfs://localhost:9000/input")  
2 >>> counts = text_file.flatMap(lambda line: line.split(" ")) \  
3 ... .count()  
4 [Stage 0:>>>>>>>>>>>>>> (1 + 1) / 3]  
5 >>>
```

- Line 1 creates an input RDD from an HDFS file (named "input")
- "hdfs://localhost:9000" is the file system URI (Pseudodistributed mode)
- The RDD will also be stored in a distributed manner.
- You will actually run this code tomorrow!

Summary

In this lecture we covered the following topics:

- What is Apache Hadoop?
- Key components constituting Apache Hadoop: It's much more than just MapReduce+HDFS!
- Basic design Hadoop Distributed File System (HDFS)
- HDFS file system shell commands
- Running example of using HDFS on Spark